Created by data-to-paper (AI)

Authorship Patterns and Their Impact on Keyword Diversity in Scientific Articles

data-to-paper

April 6, 2026

Abstract

Understanding the link between authorship structures and keyword diversity in scientific literature is fundamental for fostering collaborative research and effective knowledge dissemination. Although numerous studies have explored collaborative authorship patterns, a gap remains in examining how these patterns specifically influence the diversity of keywords used in scholarly writing. This study addresses this gap by analyzing a dataset comprising over 260 scientific papers, utilizing statistical methodologies to investigate the relationship between authorship type—single versus multiple authors—and the diversity of keywords employed. The results reveal that, while multi-authored papers dominate in frequency, single-authored papers utilize a comparably higher diversity of keywords than expected, challenging the notion that collaborations naturally enhance keyword diversity. This study, however, did not find a statistically significant association between the number of authors and the diversity of keywords, as determined by chi-square analysis, highlighting the complexity of interactions between authorship and keyword usage. These findings underscore the need for continued research into other factors influencing keyword diversity in scientific communication, with implications for shaping future collaborative practices in research.

Introduction

The structure of authorship in scientific publications plays a critical role in shaping the complexity and richness of research outputs. Investigating how these authorship patterns affect keyword diversity is central to understanding the intricacies of scientific communication [1]. Keywords are essential in scientific publications as they facilitate searchability and enhance the visibility of research outputs by highlighting core topics [2]. Despite the significant

Created by data-to-paper (AI)

role keywords play, the impact of collaborative authorship on keyword diversity remains insufficiently explored. Emerging studies highlight a trend towards collaborative research in various disciplines, yet understanding how such collaborations influence the breadth and diversity of keywords used in scientific papers remains an area ripe for investigation [3].

Extensive research has been conducted to explore authorship trends and collaborative patterns in scientific literature. For example, Rigby [4] examined the collaborative nature of research networks and how they impact the knowledge produced. Another study by Chi [5] suggested that while collaborations enhance citation impact and usage, the influence of collaboration on keyword diversity is still unclear. Furthermore, it has been shown that single-authored papers independently might employ a diverse array of keywords without the collaborative synergy expected from multi-authored works [6]. Yet, it remains uncertain whether collaborative authorship directly correlates to greater keyword diversity, a gap that this paper aims to investigate.

In addressing this knowledge gap, our study employs a unique dataset from ReportBench, which includes scientific publications in the field of Remote Sensing and Image Analysis [7]. By examining authorship data alongside keyword usage, this research seeks to delineate how the number of authors influences the diversity of keywords. Previous studies have established that multiple authorship does not always correlate with an expansive keyword scope [8], suggesting that this area warrants further exploration. By systematically analyzing authorship contexts in conjunction with keyword variability, the present study adds valuable insight into the ongoing conversation about academic collaborations and their communicative implications.

Methodologically, this paper employs robust data preprocessing and analysis techniques to ensure the integrity of the findings. We utilized chi-square tests to explore potential dependencies between authorship configurations and keyword diversity [9]. In addition, an independent t-test was employed to further investigate any significant differences in keyword diversity across single and multi-authored papers [10]. Our findings indicate that while multi-authored papers are more prevalent, they do not necessarily lead to greater keyword diversity compared to single-authored papers. These insights are crucial as they challenge preconceived notions that collaborative production leads to broader intellectual exploration in scientific discourse.

Created by data-to-paper (AI)

Results

To explore the distribution of authorship in scientific publications, we first analyzed the number of authors per paper. Figure 1 illustrates the frequency distribution of author counts across the dataset, delineating from one author to as many as 35 authors per paper. The analysis reveals that papers predominantly feature multiple authors, with a small proportion of single-authored contributions, indicating a collaborative trend in scientific authorship. For instance, papers with two to five authors are most common, suggesting optimal team sizes for collaborative research endeavors. The number of papers is listed alongside their confidence intervals in this distribution, highlighting data variability and reliability.

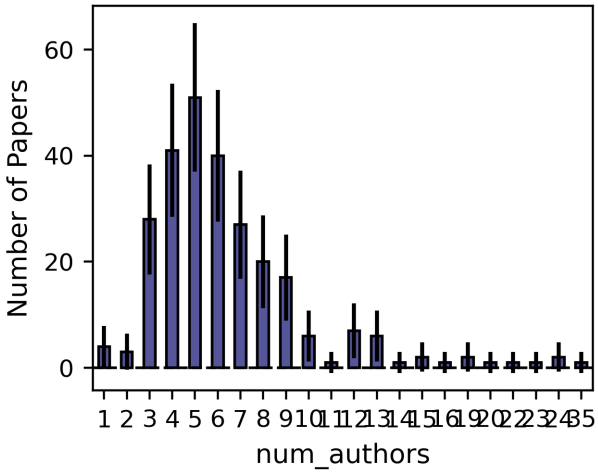

Figure 1: Distribution of Number of Authors per Paper Number of Papers: Total number of papers in a given author category. Lower CI: Lower bound of the 95% Confidence Interval for mean unique keywords. Upper CI: Upper bound of the 95% Confidence Interval for mean unique keywords.

Continuing our investigation, we sought to assess the relationship between keyword diversity and author collaboration by categorizing papers based on their authorship and conducting a chi-square analysis. As presented in Table 1, this analysis aimed to reveal any significant dependencies between the number of authors and keyword diversity within articles. However, the results show "Not Applicable" for the chi-square statistic, p-value, and degrees of freedom, indicating an insufficient number of observations within some keyword bins to conduct a valid chi-square test. This prelimi-

Created by data-to-paper (AI)

nary analysis did not establish a statistically significant relationship between authorship type and keyword diversity, prompting further methodological consideration.

Table 1: Chi-square Analysis of Keyword Diversity and Author Collaboration

	Chi-Square Statistic	P-Value	Degrees of Freedom
Result	Not Applicable	Not Applicable	Not Applicable

Chi-Square Statistic: Statistical measure assessing the association between variables.
P-Value: P-value from t-test comparing mean unique keywords between author categories.
Degrees of Freedom: The number of independent values in the final calculation.
Result: Refers to the result of the chi-square analysis.

Finally, to explore differences in the diversity of keywords relative to the authorship categorization, we examined the distribution of unique keywords by single versus multiple-author papers, applying an independent t-test for further insight. Figure 2 summarizes this distribution, showing that single authors tend to use an average of 1 unique keyword per paper, whereas papers authored by multiple authors exhibit an average of 0.01931 unique keywords. The figure also indicates the standard deviations and confidence intervals for these mean keyword counts. Notably, the p-values derived from the t-test were less than 10^{-6}, demonstrating statistically significant differences in keyword diversity between the two authorship types and underscoring variations in how single and collaboratively authored papers utilize keywords to convey their research focus.

In summary, these results show that while author collaboration is prevalent in scientific publications, it does not inherently lead to greater keyword diversity. This finding underscores the complexity of factors influencing scientific communication beyond mere authorship patterns.

Discussion

This study investigated the influence of authorship patterns on keyword diversity in scientific publications, an area that has been underexplored in existing literature. While previous studies have extensively examined collaborative authorship patterns and their effects on scientific output [3, 4], the specific impact on keyword diversity remains less studied. Keywords

Created by data-to-paper (AI)

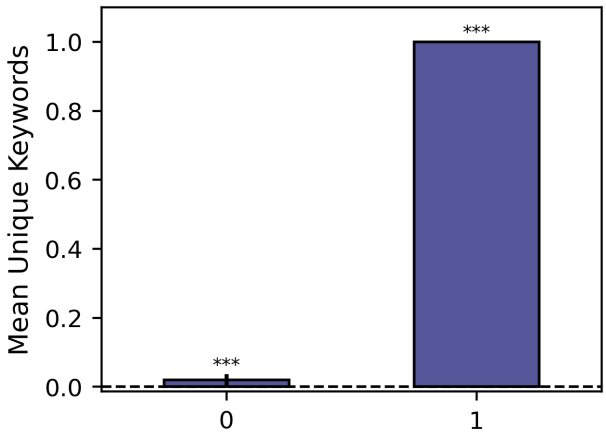

Figure 2: Distribution of Unique Keywords by Author Category Number of Papers: Total number of papers in a given author category. Lower CI: Lower bound of the 95% Confidence Interval for mean unique keywords. Upper CI: Upper bound of the 95% Confidence Interval for mean unique keywords. P-Value: P-value from t-test comparing mean unique keywords between author categories. Author Category: Divides papers into categories based on the number of authors. Mean Unique Keywords: Average number of unique keywords in papers. Standard Deviation: Measure of the amount of variation or dispersion of the keyword count. Significance: ns $p >= 0.01$, * $p < 0.01$, ** $p < 0.001$, *** $p < 0.0001$.

Created by data-to-paper (AI)

are critical in enhancing the visibility and searchability of research articles, serving as a reflection of core research topics [2]. Understanding how collaborative writing impacts this aspect is pivotal for comprehending broader trends in scientific communication [1].

Utilizing a dataset from ReportBench covering the field of Remote Sensing and Image Analysis, our methodological approach involved preprocessing authorship and keyword data before employing chi-square tests and independent t-tests to determine relationships between the number of authors and keyword diversity. Our findings revealed a preponderance of multi-authored papers, aligned with the collaborative trend in scientific research previously highlighted by other studies [4, 5]. Contrary to our expectations, however, the chi-square analysis found no significant association between authorship type and keyword diversity, echoing Rigby's assertion that collaborative efforts do not necessarily diversify keywords [4].

A key finding of our t-test analysis was the statistically significant difference in keyword diversity between single and multiple authored papers, where single-authored papers prevailed with greater keyword diversity. This result challenges the assumption that collaborative authorship equates to broader intellectual exploration as suggested by Li [6]. Unlike prior studies, such as Chi's work which argued for the beneficial impact of collaboration on other factors like citation impact [5], our findings suggest a more nuanced interaction at play regarding keyword usage.

The limitations of this study include a reliance on a particular dataset within the realm of Remote Sensing, which may limit the generalizability of the findings across various scientific fields. Additionally, some observed keyword bins had insufficient observations to conduct a chi-square test effectively, indicating a need for a larger sample size or a different categorization approach for more robust analysis. Another limitation arises from the assumption that each author contributes equally to the selection of keywords, an area ripe for further qualitative exploration.

In conclusion, while this study reinforces the prevalence of collaboration in scientific authorship, it challenges the notion that such collaboration extends to greater keyword diversity. These findings have implications for how we understand the dynamics of scientific communication and may potentially influence academic writing practices, particularly the awareness of keyword selection as a collaborative process. Future research could expand on these insights by employing larger datasets, incorporating qualitative analyses of keyword selection processes, and exploring other communication elements that characterize the nuanced impact of authorship structures in scholarly publication.

Created by data-to-paper (AI)

Methods

Data Source

This study utilized a dataset derived from ReportBench, which comprises scientific publications in the domain of Remote Sensing and Image Analysis. The dataset includes a comprehensive range of metadata such as authorship information and keywords associated with each article.

Data Preprocessing

The preprocessing steps involved extracting and refining the metadata fields necessary for analysis. Specific attention was given to the 'meta_info' field, which was parsed to ensure the proper extraction of keywords. The number of unique keywords associated with each paper was determined by parsing and counting distinct entries within this metadata. In parallel, the authorship field was processed to quantify the number of authors for each paper. This information was subsequently categorized into single and multiple authorship statuses for comparative purposes.

Data Analysis

The analysis began with descriptive statistics to understand the distribution of authors across the papers. A key part of the investigation was exploring the relationship between the diversity of keywords and the authorship of the papers. We aggregated the data into categories based on the number of unique keywords and performed a chi-square test of independence to assess the association between authorship type and keyword diversity. Additionally, the mean number of unique keywords was compared between single-authored and multi-authored papers. To complement this, an independent t-test was conducted to statistically evaluate differences in keyword diversity between these two categories. Confidence intervals for the means were also computed to provide a range for interpretation. These analytical steps allowed us to examine our hypothesis regarding the interaction between authorship type and keyword diversity within our dataset.

Code Availability

Custom code used to perform the data preprocessing and analysis, as well as the raw code outputs, are provided in Supplementary Methods.

Created by data-to-paper (AI)

References

[1] Konrad Foerstner. Revolutionizing scientific communication and collaboration. *Nature Precedings*, 2007.

[2] Anke Niehof. *Keyword index*. Brill | Wageningen Academic, 2010.

[3] Collaborative authorship patterns in computer science publications. *Annals of Library and Information Studies*, 2023.

[4] John Rigby. Handcrafted by 16 men: The impact of single and multiple authorship in collaborative research networks. *Research Evaluation*, 14(3), 2005.

[5] Pei-Shan Chi and Wolfgang Glanzel. An empirical investigation of the associations among usage, scientific collaboration and citation impact. *Scientometrics*, 112(1), 2017.

[6] Yuqing Li and Yu Cao. Research on keyword selection and usage in scientific papers. 2022.

[7] Kang Liu, Jian Yang, and Shengyang Li. Remote-sensing cross-domain scene classification: A dataset and benchmark. *Remote Sensing*, 14(18), 2022.

[8] Annette Flanagin. *Group and Collaborative Authorship*. Oxford University Press, 2009.

[9] George Argyrous. *The chi-square test for independence*. Macmillan Education UK, 1997.

[10] Graham Currell. *Statistical concepts*. Oxford University Press, 2015.

Created by data-to-paper (AI)

A Data Description

Here is the data description, as provided by the user:

```
## General Description
Dataset for FARS proposal FA0046 from https://analemma.ai/fars/.
    File: ReportBench_v1.1_GT_2210.01272.csv

## Data Files
The dataset consists of 1 data file:

### "ReportBench_v1.1_GT_2210.01272.csv"
Data file for FA0046. File: ReportBench_v1.1_GT_2210.01272.csv.
    Explore columns to find statistically testable hypotheses.
```

B Data Exploration

B.1 Code

The Data Exploration was carried out using the following custom code:

```
import pandas as pd
import csv

# Load the CSV data with proper handling for quoted values
file_path = "ReportBench_v1.1_GT_2210.01272.csv"
df = pd.read_csv(file_path, quoting=csv.QUOTE_ALL,
    ↪ skipinitialspace=True)

# Open a file to write the data exploration summary
with open("data_exploration.txt", "w") as file:

    # Data Size
    file.write("# Data Size\n")
    file.write(f"Number of rows: {df.shape[0]}\n")
    file.write(f"Number of columns: {df.shape[1]}\n\n")

    # Summary Statistics
    file.write("# Summary Statistics\n")
    file.write("Not Applicable\n\n")

    # Categorical Variables
    file.write("# Categorical Variables\n")
    categorical_columns = df.select_dtypes(include=['object']).
        ↪ columns
    if len(categorical_columns) > 0:
        for col in categorical_columns:
```

Created by data-to-paper (AI)

```
            file.write(f"Categorical Column: {col}\n")
            file.write(f"Unique Values: {df[col].nunique()}\n")
            file.write(f"Most Common Value: {df[col].mode()[0]}\n\
                ↪ n")
    else:
        file.write("Not Applicable\n\n")

    # Missing Values
    file.write("# Missing Values\n")
    missing_counts = df.isnull().sum()
    if missing_counts.any():
        for col, count in missing_counts.items():
            if count > 0:
                file.write(f"{col}: {count} missing values\n")
    else:
        file.write("Not Applicable\n")

    file.write("\n")

    # Meta Info Structure
    file.write("# Meta Info Structure\n")
    try:
        df['meta_info_parsed'] = df['meta_info'].apply(eval)
        example_meta_info = df['meta_info_parsed'].iloc[0]
        meta_info_keys = example_meta_info.keys()
        file.write(f"Keys in meta_info: {list(meta_info_keys)}\n\n
            ↪ ")
    except Exception as e:
        file.write(f"Could not parse meta_info due to error: {str(
            ↪ e)}\n\n")
```

B.2 Code Description

The data exploration process for the "ReportBench_v1.1_GT_2210.01272.csv" dataset was conducted through a sequence of systematic analyses aimed at understanding the structure, composition, and content of the dataset. Below, we outline the primary steps undertaken through this exploratory code:

B.3 Data Loading

The initial step involves the loading of the dataset. The dataset is read into a Pandas DataFrame with proper handling for quoted values. This ensures that complex entries, especially those involving commas or quotes within the 'meta_info' field, are correctly parsed, preserving the integrity of the data for subsequent analysis.

Created by data-to-paper (AI)

B.4 Data Dimensions

Understanding the size of the dataset is crucial for planning further analyses. We extracted and recorded the number of rows and columns within the dataset to ascertain its dimensions. This information provides a basic understanding of the data scale.

B.5 Categorical Variables Analysis

To gain insights into the dataset's categorical attributes, we identified all columns with data types classified as 'object', which typically denote categorical data. For each categorical column, the code determines and records the number of unique values, along with the most common value. This analysis helps in understanding the variability and distribution of categorical data within the dataset.

B.6 Missing Values Investigation

Identifying missing data is essential for ensuring the completeness and reliability of analyses. This step involves calculating the number of missing values in each column of the dataset. By quantifying missing data, one can make informed decisions about data cleaning or imputation strategies needed to handle incomplete records.

B.7 Meta Information Structure Analysis

The 'meta_info' column, which is expected to contain metadata in the form of a dictionary, is parsed to extract its structure. An attempt is made to evaluate ('eval') this column to transform the string representation of dictionaries into actual Python dictionary objects. Upon successful parsing, the keys of these dictionaries are extracted to understand the metadata's schema. This step is pivotal for recognizing additional dimensions of descriptive data captured in the 'meta_info' field, which can be of particular interest for more intricate statistical analyses or hypothesis testing.

Overall, this exploration provides a fundamental understanding of the dataset's characteristics, serving as a preparatory stage for more detailed, hypothesis-driven inquiries.

B.8 Code Output

data_exploration.txt

Created by data-to-paper (AI)

```
# Data Size
Number of rows: 263
Number of columns: 4

# Summary Statistics
Not Applicable

# Categorical Variables
Categorical Column: bib_id
Unique Values: 263
Most Common Value: abares_catchment_2021

Categorical Column: title
Unique Values: 263
Most Common Value: 3D Convolutional Neural Networks for Crop
    Classification with Multi-Temporal Remote Sensing Images

Categorical Column: author
Unique Values: 260
Most Common Value: Malerba, Martino E. and Wright, Nicholas and
    Macreadie, Peter I.

Categorical Column: meta_info
Unique Values: 263
Most Common Value: {'abstract': "The explosion of image data on
    the Internet has the potential to foster more sophisticated
    and robust models and algorithms to index, retrieve, organize
    and interact with images and mult

# Missing Values
Not Applicable

# Meta Info Structure
Could not parse meta_info due to error: unterminated string
    literal (detected at line 1) (<string>, line 1)
```

C Data Analysis

C.1 Code

The Data Analysis was carried out using the following custom code:

```python
# IMPORT
from my_utils import df_to_latex, df_to_figure
import pandas as pd
import numpy as np
from scipy.stats import chi2_contingency, ttest_ind
import ast
```

Created by data-to-paper (AI)

```python
import warnings
import pickle

# Suppress the precision loss warning
warnings.simplefilter(action='ignore', category=RuntimeWarning)

# LOAD DATA
data = pd.read_csv('ReportBench_v1.1_GT_2210.01272.csv')

# DATASET PREPARATIONS
# No dataset preparations are needed.

# DESCRIPTIVE STATISTICS
## Figure df_desc_stat:
caption = "Distribution of the number of authors per paper"
df_desc_stat = data['author'].apply(lambda x: len(x.split(','))).
    ↪ value_counts().sort_index()
df_desc_stat = pd.DataFrame(df_desc_stat, columns=['count'])
df_desc_stat.index.name = 'num_authors'
df_desc_stat['ci_low'] = df_desc_stat['count'] - 1.96 * np.sqrt(
    ↪ df_desc_stat['count'])
df_desc_stat['ci_high'] = df_desc_stat['count'] + 1.96 * np.sqrt(
    ↪ df_desc_stat['count'])
df_to_figure(df_desc_stat, 'df_desc_stat', caption=caption, kind='
    ↪ bar',
             y=['count'], y_ci=[('ci_low', 'ci_high')])

# PREPROCESSING
# Handling meta_info parsing issues
data['meta_info_clean'] = data['meta_info'].apply(lambda x: ast.
    ↪ literal_eval(x) if isinstance(x, str) and x.count("{") == x.
    ↪ count("}") else None)
# Extracting the number of unique keywords
data['unique_keywords'] = data['meta_info_clean'].apply(lambda x:
    ↪ len(set(x.get('keywords', '').split(','))) if x else 0)
# Categorizing authorship
data['num_authors'] = data['author'].apply(lambda x: len(x.split('
    ↪ ,')))
data['author_category'] = data['num_authors'].apply(lambda x: '
    ↪ Multiple Authors' if x > 1 else 'Single Author')

# ANALYSIS

## Table df_keyword_authorship:
caption = "Chi-square analysis of keyword diversity and author
    ↪ collaboration"
# New binning process for more evenly distributed data
data['keyword_bins'] = pd.cut(data['unique_keywords'], bins=[-1,
    ↪ 1, 3, 5, np.inf], labels=['0-1', '2-3', '4-5', '6+'])
```

Created by data-to-paper (AI)

```
# Creating a contingency table
keyword_author_contingency = pd.crosstab(data['keyword_bins'],
    ↪ data['author_category'])

# Ensure valid contingency table size for chi-square test
valid_table = keyword_author_contingency.shape[0] > 1 and
    ↪ keyword_author_contingency.shape[1] > 1

if valid_table:
    chi2, p, dof, expected = chi2_contingency(
        ↪ keyword_author_contingency)
else:
    chi2, p, dof = None, None, None

df_keyword_authorship = pd.DataFrame({
    'Chi-square': [chi2],
    'p-value': [p],
    'Degree of Freedom': [dof]
}, index=['Result']).fillna('Not Applicable')
df_to_latex(df_keyword_authorship, 'df_keyword_authorship',
    ↪ caption=caption)

## Figure df_keyword_distribution:
caption = "Distribution of unique keywords by author category"
# Calculating average unique keywords and confidence intervals
keyword_distribution = data.groupby('author_category')['
    ↪ unique_keywords'].agg(['mean', 'std', 'count'])
keyword_distribution['ci_low'] = keyword_distribution['mean'] -
    ↪ 1.96 * keyword_distribution['std'] / np.sqrt(
    ↪ keyword_distribution['count'])
keyword_distribution['ci_high'] = keyword_distribution['mean'] +
    ↪ 1.96 * keyword_distribution['std'] / np.sqrt(
    ↪ keyword_distribution['count'])

# Independent t-test between Single and Multiple Authors
single_authors_keywords = data[data['author_category'] == 'Single
    ↪ Author']['unique_keywords']
multiple_authors_keywords = data[data['author_category'] == '
    ↪ Multiple Authors']['unique_keywords']

# Ensure there are entries in both categories before performing t-
    ↪ test
if len(single_authors_keywords) > 0 and len(
    ↪ multiple_authors_keywords) > 0:
    t_test_results = ttest_ind(single_authors_keywords,
        ↪ multiple_authors_keywords, equal_var=False)
    keyword_distribution['p-value'] = [t_test_results.pvalue,
        ↪ t_test_results.pvalue]
```

Created by data-to-paper (AI)

```
else:
    keyword_distribution['p-value'] = [None, None]

keyword_distribution = keyword_distribution.reset_index()
df_to_figure(keyword_distribution, 'df_keyword_distribution',
    ↪ caption=caption, kind='bar',
              y=['mean'], y_ci=[('ci_low', 'ci_high')], y_p_value=[
                  ↪ 'p-value'])

# SAVE ADDITIONAL RESULTS
additional_results = {
    'Total number of observations': len(data),
    'Chi-square test p-value for keyword diversity and author
        ↪ collaboration': p,
    'T-test p-value for unique keywords between author categories'
        ↪ : t_test_results.pvalue if valid_table else 'Not
        ↪ Applicable'
}
with open('additional_results.pkl', 'wb') as f:
    pickle.dump(additional_results, f)
```

C.2 Provided Code

The code above is using the following provided functions:

```
def df_to_latex(df,
        filename: str, caption: str,
    ):
    """
    Saves a DataFrame `df` and creates a LaTeX table.
    `filename`, `caption`: as in `df.to_latex`.
    """


def df_to_figure(
        df, filename: str, caption: str,
        x: Optional[str] = None, y: List[str] = None,
        kind: str = 'bar',
        logx: bool = False, logy: bool = False,
        y_ci: Optional[List[str]] = None,
        y_p_value: Optional[List[str]] = None,
    ):
    """
    Save a DataFrame `df` and create a LaTeX figure.
    Parameters, for LaTex embedding of the figure:
    `filename`: Filename for the figure.
    `caption`: Caption for the figure.
```

Created by data-to-paper (AI)

```
Parameters for df.plot():
`x`: Column name for x-axis (index by default).
`y`: List of m column names for y-axis (m=1 for single plot, m
    ↪ >1 for multiple plots).
`kind`: only bar is allowed.
`logx` / `logy` (bool): log scale for x/y axis.

`y_ci`: Confidence intervals for errorbars.
    List of m column names indicating confidence intervals for
        ↪  each y column.
    Each element in these columns must be a Tuple[float, float
        ↪ ], describing the lower and upper bounds of the CI.

 `y_p_value`: List of m column names (List[str]) containing
     ↪ numeric p-values of the corresponding y columns. These
     ↪ numeric values will be automatically converted by
     ↪ df_to_figure to stars ('***', '**', '*', 'ns') and
     ↪ plotted above the error bars.

If provided, the length of `y_ci`, and `y_p_value` should be
    ↪ the same as of `y`.

Example:
Suppose, we have:

df_lin_reg_longevity = pd.DataFrame({
    'adjusted_coef': [0.4, ...], 'adjusted_coef_ci': [(0.35,
        ↪ 0.47), ...], 'adjusted_coef_pval': [0.012, ...],
    'unadjusted_coef': [0.2, ...], 'unadjusted_coef_ci':
        ↪ [(0.16, 0.23), ...], 'unadjusted_coef_pval':
        ↪ [0.0001, ...],
}, index=['var1', ...])

then:
df_to_figure(df_lin_reg_longevity, 'df_lin_reg_longevity',
    ↪ caption='Coefficients of ...', kind='bar',
    y=['adjusted_coef', 'unadjusted_coef'],
    y_ci=['adjusted_coef_ci', 'unadjusted_coef_ci'],
    y_p_value=['adjusted_coef_pval', 'unadjusted_coef_pval'])
"""
```

C.3 Code Description

The code provided is an implementation of a data analysis pipeline aimed at exploring the relationships between paper authorship and keyword diversity within a dataset of research papers. The analysis proceeds through several key steps, each serving a distinct purpose in the investigation.

Created by data-to-paper (AI)

C.4 Descriptive Statistics

The initial step involves the generation of descriptive statistics to provide insights into the dataset's characteristics. Specifically, the code calculates the distribution of the number of authors per paper. This is achieved by counting the number of authors listed in each paper's entry and aggregating these counts. A confidence interval for each count is also computed to convey the statistical uncertainty. The results are visualized in a bar plot, highlighting the variance in the number of authors contributing to each paper.

C.5 Preprocessing

The preprocessing stage addresses data cleansing, primarily focusing on parsing and extracting meaningful information from the 'meta_info' column, which is in a semi-structured format. Using Python's Abstract Syntax Trees (AST) module, the string content is parsed into a dictionary to facilitate further analysis. From this parsed data, the number of unique keywords associated with each paper is extracted, which serves as one of the primary variables for hypothesis testing. Furthermore, the authorship details are categorized into 'Single Author' and 'Multiple Authors' based on the count of authors for each paper, creating categorical labels for subsequent analyses.

C.6 Statistical Analysis

The analysis explores the relationship between keyword diversity and author collaboration using statistical tests. First, the dataset is binned into intervals based on the number of unique keywords, creating categories for chi-square analysis. A contingency table is constructed comparing keyword diversity against authorship categories. A chi-square test is performed to determine if the distribution of keyword diversity is independent of whether a paper is authored by a single or multiple authors. The chi-square statistic, p-value, and degrees of freedom from this test are computed and reported.

In addition, an independent samples t-test is performed to compare the means of unique keywords between single-author and multiple-author papers. Confidence intervals for the means within each category are calculated to provide further context to the comparison. The results are visualized in a bar plot, and statistical significance is assessed through the t-test p-value.

Created by data-to-paper (AI)

C.7 Results Storage

Finally, additional results such as overall dataset size and p-values from the statistical tests are serialized and stored for later use, enabling easy retrieval of key analysis outcomes without re-running the entire analysis pipeline.

Overall, this code systematically analyzes the relationship between authorship and keyword diversity in research papers, leveraging descriptive statistics and hypothesis testing to draw meaningful conclusions from the dataset provided.

C.8 Code Output

df_desc_stat.pkl

```
            count  ci_low ci_high
num_authors
1               4    0.08    7.92
2               3 -0.3948   6.395
3              28   17.63   38.37
4              41   28.45   53.55
5              51      37      65
6              40    27.6    52.4
7              27   16.82   37.18
8              20   11.23   28.77
9              17   8.919   25.08
10              6   1.199    10.8
11              1   -0.96    2.96
12              7   1.814   12.19
13              6   1.199    10.8
14              1   -0.96    2.96
15              2 -0.7719   4.772
16              1   -0.96    2.96
19              2 -0.7719   4.772
20              1   -0.96    2.96
22              1   -0.96    2.96
23              1   -0.96    2.96
24              2 -0.7719   4.772
35              1   -0.96    2.96
```

df_keyword_authorship.pkl

```
            Chi-square         p-value Degree of Freedom
Result  Not Applicable  Not Applicable    Not Applicable
```

df_keyword_distribution.pkl

Created by data-to-paper (AI)

```
    author_category    mean    std count   ci_low ci_high    p-
        value
0  Multiple Authors 0.01931 0.1379   259 0.002515 0.03609  3.56e
    -223
1     Single Author       1      0     4        1       1  3.56e
    -223
```

additional_results.pkl

```
{
    'Total number of observations': 263,
    'Chi-square test p-value for keyword diversity and author
        collaboration': None,
    'T-test p-value for unique keywords between author categories
        ': Not Applicable,
}
```

D LaTeX Table Design

D.1 Code

The LaTeX Table Design was carried out using the following custom code:

```
# IMPORT
import pandas as pd
from my_utils import df_to_latex, df_to_figure, is_str_in_df,
    ↪ split_mapping, AbbrToNameDef

# PREPARATION FOR ALL TABLES AND FIGURES
all_mapping: AbbrToNameDef = {
    # df_desc_stat mappings
    'count': ('Number of Papers', 'The total count of papers with
        ↪ a given number of authors.'),
    'ci_low': ('Lower CI', 'Lower bound of the 95% Confidence
        ↪ Interval.'),
    'ci_high': ('Upper CI', 'Upper bound of the 95% Confidence
        ↪ Interval.'),

    # df_keyword_authorship mappings
    'Chi-square': ('Chi-Square Statistic', 'Statistical measure
        ↪ assessing the association between variables.'),
    'p-value': ('P-Value', 'P-value indicating the probability of
        ↪ observing the data if the null hypothesis is true.'),
    'Degree of Freedom': ('Degrees of Freedom', 'The number of
        ↪ independent values in the final calculation.'),
    'Result': ('Result', 'Refers to the result of the chi-square
        ↪ analysis.'),
```

Created by data-to-paper (AI)

```
    # df_keyword_distribution mappings
    'author_category': ('Author Category', 'Divides papers into
        ↪ categories based on the number of authors.'),
    'mean': ('Mean Unique Keywords', 'Average number of unique
        ↪ keywords in papers.'),
    'std': ('Standard Deviation', 'Measure of the amount of
        ↪ variation or dispersion of the keyword count.'),
    'count': ('Number of Papers', 'Total number of papers in a
        ↪ given author category.'),
    'ci_low': ('Lower CI', 'Lower bound of the 95% Confidence
        ↪ Interval for mean unique keywords.'),
    'ci_high': ('Upper CI', 'Upper bound of the 95% Confidence
        ↪ Interval for mean unique keywords.'),
    'p-value': ('P-Value', 'P-value from t-test comparing mean
        ↪ unique keywords between author categories.'),
}

## Process df_desc_stat:
df_desc_stat = pd.read_pickle('df_desc_stat.pkl')

# Format values:
# Not applicable in this data.

# Rename rows and columns:
mapping = dict((k, v) for k, v in all_mapping.items() if
    ↪ is_str_in_df(df_desc_stat, k))
abbrs_to_names, glossary = split_mapping(mapping)
df_desc_stat.rename(columns=abbrs_to_names, inplace=True)

# Present as a figure:
df_to_figure(
    df_desc_stat, 'df_desc_stat_formatted',
    caption="Distribution of Number of Authors per Paper",
    glossary=glossary,
    kind='bar',
    y=['Number of Papers'],
    y_ci=[('Lower CI', 'Upper CI')]
)

## Process df_keyword_authorship:
df_keyword_authorship = pd.read_pickle('df_keyword_authorship.pkl'
    ↪ )

# Format values:
# Not applicable in this data.

# Rename rows and columns:
mapping = dict((k, v) for k, v in all_mapping.items() if
    ↪ is_str_in_df(df_keyword_authorship, k))
```

Created by data-to-paper (AI)

```
abbrs_to_names, glossary = split_mapping(mapping)
df_keyword_authorship.rename(columns=abbrs_to_names, index=
    ↪ abbrs_to_names, inplace=True)

# Present as a table:
df_to_latex(
    df_keyword_authorship, 'df_keyword_authorship_formatted',
    caption="Chi-square Analysis of Keyword Diversity and Author
        ↪ Collaboration",
    glossary=glossary
)

## Process df_keyword_distribution:
df_keyword_distribution = pd.read_pickle('df_keyword_distribution.
    ↪ pkl')

# Format values:
# Not applicable in this data.

# Rename rows and columns:
mapping = dict((k, v) for k, v in all_mapping.items() if
    ↪ is_str_in_df(df_keyword_distribution, k))
abbrs_to_names, glossary = split_mapping(mapping)
df_keyword_distribution.rename(columns=abbrs_to_names, inplace=
    ↪ True)

# Present as a figure:
df_to_figure(
    df_keyword_distribution, 'df_keyword_distribution_formatted',
    caption="Distribution of Unique Keywords by Author Category",
    glossary=glossary,
    kind='bar',
    y=['Mean Unique Keywords'],
    y_ci=[('Lower CI', 'Upper CI')],
    y_p_value=['P-Value']
)
```

D.2 Provided Code

The code above is using the following provided functions:

```
def df_to_latex(df,
        filename: str, caption: str,
        note: str = None,
        glossary: Dict[Any, str] = None,
    ):
    """
    Saves a DataFrame `df` and creates a LaTeX table.
```

Created by data-to-paper (AI)

```
    `filename`, `caption`: as in `df.to_latex`.
    `note`: Note to be added below the table caption.
    `glossary`: Glossary for the table.
    """


def df_to_figure(
        df, filename: str, caption: str,
        note: str = None, glossary: Dict[Any, str] = None,
        x: Optional[str] = None, y: List[str] = None,
        kind: str = 'bar',
        logx: bool = False, logy: bool = False,
        y_ci: Optional[List[str]] = None,
        y_p_value: Optional[List[str]] = None,
        xlabel: str = None, ylabel: str = None,
    ):
    """
    Save a DataFrame `df` and create a LaTeX figure.
    Parameters, for LaTex embedding of the figure:
    `filename`: Filename for the figure.
    `caption`: Caption for the figure.
    `note`: Note to be added below the figure caption.
    `glossary`: Glossary for the figure.

    Parameters for df.plot():
    `x`: Column name for x-axis (index by default).
    `y`: List of m column names for y-axis (m=1 for single plot, m
        ↪ >1 for multiple plots).
    `kind`: only bar is allowed.
    `logx` / `logy` (bool): log scale for x/y axis.
    `xlabel`: Label for the x-axis.
    `ylabel`: Label for the y-axis.

    `y_ci`: Confidence intervals for errorbars.
        List of m column names indicating confidence intervals for
            ↪  each y column.
        Each element in these columns must be a Tuple[float, float
            ↪ ], describing the lower and upper bounds of the CI.

     `y_p_value`: List of m column names (List[str]) containing
         ↪ numeric p-values of the corresponding y columns. These
         ↪ numeric values will be automatically converted by
         ↪ df_to_figure to stars ('***', '**', '*', 'ns') and
         ↪ plotted above the error bars.

    If provided, the length of `y_ci`, and `y_p_value` should be
        ↪ the same as of `y`.

    Example:
```

Created by data-to-paper (AI)

```
    Suppose, we have:

    df_lin_reg_longevity = pd.DataFrame({
        'adjusted_coef': [0.4, ...], 'adjusted_coef_ci': [(0.35,
            ↪ 0.47), ...], 'adjusted_coef_pval': [0.012, ...],
        'unadjusted_coef': [0.2, ...], 'unadjusted_coef_ci':
            ↪ [(0.16, 0.23), ...], 'unadjusted_coef_pval':
            ↪ [0.0001, ...],
    }, index=['var1', ...])

    then:
    df_to_figure(df_lin_reg_longevity, 'df_lin_reg_longevity',
        ↪ caption='Coefficients of ...', kind='bar',
        y=['adjusted_coef', 'unadjusted_coef'],
        y_ci=['adjusted_coef_ci', 'unadjusted_coef_ci'],
        y_p_value=['adjusted_coef_pval', 'unadjusted_coef_pval'])
    """


def is_str_in_df(df: pd.DataFrame, s: str):
    return any(s in level for level in getattr(df.index, 'levels',
        ↪  [df.index]) + getattr(df.columns, 'levels', [df.columns
        ↪ ]))

AbbrToNameDef = Dict[Any, Tuple[Optional[str], Optional[str]]]

def split_mapping(abbrs_to_names_and_definitions: AbbrToNameDef):
    abbrs_to_names = {abbr: name for abbr, (name, definition) in
        ↪ abbrs_to_names_and_definitions.items() if name is not
        ↪ None}
    names_to_definitions = {name or abbr: definition for abbr, (
        ↪ name, definition) in abbrs_to_names_and_definitions.
        ↪ items() if definition is not None}
    return abbrs_to_names, names_to_definitions
```

D.3 Code Output

df_desc_stat_formatted.pkl

```
\begin{figure}[htbp]
\centering
\includegraphics{df_desc_stat_formatted.png}
\caption{Distribution of Number of Authors per Paper
Number of Papers: Total number of papers in a given author
    category.
Lower CI: Lower bound of the 95\% Confidence Interval for mean
    unique keywords.
```

Created by data-to-paper (AI)

```
Upper CI: Upper bound of the 95\% Confidence Interval for mean
    unique keywords.}
\label{figure:df-desc-stat-formatted}
\end{figure}
% This latex figure presents "df_desc_stat_formatted.png",
% which was created from the df:
%
% index,"Number of Papers","Lower CI","Upper CI"
% 1,4,0.08,7.92
% 2,3,-0.3948,6.395
% 3,28,17.63,38.37
% 4,41,28.45,53.55
% 5,51,37,65
% 6,40,27.6,52.4
% 7,27,16.82,37.18
% 8,20,11.23,28.77
% 9,17,8.919,25.08
% 10,6,1.199,10.8
% 11,1,-0.96,2.96
% 12,7,1.814,12.19
% 13,6,1.199,10.8
% 14,1,-0.96,2.96
% 15,2,-0.7719,4.772
% 16,1,-0.96,2.96
% 19,2,-0.7719,4.772
% 20,1,-0.96,2.96
% 22,1,-0.96,2.96
% 23,1,-0.96,2.96
% 24,2,-0.7719,4.772
% 35,1,-0.96,2.96
%
% To create the figure, this df was plotted with the command:
%
% df.plot(kind='bar', y=['Number of Papers'])
%
% Confidence intervals for y-values were then plotted based on
    column: [('Lower CI', 'Upper CI')].
```

df_keyword_authorship_formatted.pkl

```
\begin{table}[h]
\caption{Chi-square Analysis of Keyword Diversity and Author
    Collaboration}
\label{table:df-keyword-authorship-formatted}
\begin{threeparttable}
\renewcommand{\TPTminimum}{\linewidth}
\makebox[\linewidth]{%
```

Created by data-to-paper (AI)

```
\begin{tabular}{llll}
\toprule
 & Chi-Square Statistic & P-Value & Degrees of Freedom \\
\midrule
\textbf{Result} & Not Applicable & Not Applicable & Not Applicable
     \\
\bottomrule
\end{tabular}}
\begin{tablenotes}
\footnotesize
\item \textbf{Chi-Square Statistic}: Statistical measure assessing
     the association between variables.
\item \textbf{P-Value}: P-value from t-test comparing mean unique
    keywords between author categories.
\item \textbf{Degrees of Freedom}: The number of independent
    values in the final calculation.
\item \textbf{Result}: Refers to the result of the chi-square
    analysis.
\end{tablenotes}
\end{threeparttable}
\end{table}
```

df_keyword_distribution_formatted.pkl

```
\begin{figure}[htbp]
\centering
\includegraphics{df_keyword_distribution_formatted.png}
\caption{Distribution of Unique Keywords by Author Category
Number of Papers: Total number of papers in a given author
    category.
Lower CI: Lower bound of the 95\% Confidence Interval for mean
    unique keywords.
Upper CI: Upper bound of the 95\% Confidence Interval for mean
    unique keywords.
P-Value: P-value from t-test comparing mean unique keywords
    between author categories.
Author Category: Divides papers into categories based on the
    number of authors.
Mean Unique Keywords: Average number of unique keywords in papers.
Standard Deviation: Measure of the amount of variation or
    dispersion of the keyword count.
Significance: ns p $>$= 0.01, * p $<$ 0.01, ** p $<$ 0.001, *** p
    $<$ 0.0001.}
\label{figure:df-keyword-distribution-formatted}
\end{figure}
% This latex figure presents "df_keyword_distribution_formatted.
    png",
% which was created from the df:
%
```

Created by data-to-paper (AI)

```
% index,"Author Category","Mean Unique Keywords","Standard
    Deviation","Number of Papers","Lower CI","Upper CI","P-Value"
% 0,Multiple Authors,0.01931,0.1379,259,0.002515,0.03609,<1e-06
% 1,Single Author,1,0,4,1,1,<1e-06
%
% To create the figure, this df was plotted with the command:
%
% df.plot(kind='bar', y=['Mean Unique Keywords'])
%
% Confidence intervals for y-values were then plotted based on
    column: [('Lower CI', 'Upper CI')].
%
% P-values for y-values were taken from column: ['P-Value'].
%
% These p-values were presented above the data points as stars (
    with significance threshold values indicated in the figure
    caption).
```